

Bayesian Evals

December 2024

Evaluating AI Systems is hard

- Training set contamination

Evaluating AI Systems is hard

- Training set contamination
- Differences in formatting/prompting lead to changes in performance

Evaluating AI Systems is hard

- Training set contamination
- Differences in formatting/prompting lead to changes in performance
- *So many eval datasets to use, so little time— e.g. big bench has 204 tasks*

Evaluating AI Systems is hard

- Training set contamination
- Differences in formatting/prompting lead to changes in performance
- *So many eval datasets to use, so little time— e.g. big bench has 204 tasks*
- Can try model-generated evals but then you get an ouroboros of evals

Evaluating AI Systems is hard

- Training set contamination
- Differences in formatting/prompting lead to changes in performance
- So many eval datasets to use, so little time— e.g. big bench has 204 tasks
- Can try model-generated evals but then you get an ouroboros of evals
- How best to model uncertainty in evals?

Evaluating AI Systems is hard

- Training set contamination
- Differences in formatting/prompting lead to changes in performance
- So many eval datasets to use, so little time— e.g. big bench has 204 tasks
- Can try model-generated evals but then you get an ouroboros of evals
- How best to model uncertainty in evals?
- How to test for subjective quality of generations (not MCQA)

Evaluating AI Systems is hard

- Training set contamination
- Differences in formatting/prompting lead to changes in performance
- *So many eval datasets to use, so little time— e.g. big bench has 204 tasks*
- Can try model-generated evals but then you get an ouroboros of evals
- *How best to model uncertainty in evals?*
- How to test for subjective quality of generations (not MCQA)
 - A/B tests (e.g. chatbot arena -> Elo scores) – expensive, difficult to model uncertainty
 - Red teaming – difficult, not standardised

OLMES (Open Language Model Evaluation Standard) - Gu et al.

- Considers the uncertainty between different experiment setups
- But not the uncertainty inside each eval experiment
- (that is, inter-model uncertainty but not intra-model uncertainty)

Table 1: Scores reported in different references for LLM performances on ARC-CHALLENGE and OPENBOOKQA. Scores indicated with [†] are using multiple-choice formulation (MCF) rather than “cloze” formulation (CF) (see Section 2.1 for definitions). Entries with “?” denote either undocumented or mixed approaches across models. Different references use different evaluation setups, some of which are not fully specified, so conclusions about which models perform best are not reproducible.

Model↓	ARC-CHALLENGE Evaluations:							OPENBOOKQA Evaluations:					
	Ref1	Ref2	Ref3	Ref4	Ref5	Ref6	OLMES	Ref2	Ref4	Ref5	Ref7	Ref8	OLMES
MPT-7B	47.7	42.6			46.5		45.7	51.4	48.6				52.4
RPJ-INCITE-7B	46.3				42.8		45.3		49.4				49.0
Falcon-7B	47.9	42.4		44.5	47.5		49.7	51.6	44.6	53.0		26.0 [†]	55.2
Mistral-7B	60.0		55.5	54.9			78.6 [†]				52.2	77.6 [†]	80.6 [†]
Llama2-7B	53.1	45.9	43.2	45.9	48.5	53.7 [†]	54.2	58.6	58.6	48.4	58.6	54.4 [†]	57.8
Llama2-13B	59.4	49.4	48.8	49.4		67.6 [†]	67.3 [†]	57.0	57.0		57.0	63.4 [†]	65.4 [†]
Llama3-8B	60.2					78.6 [†]	79.3 [†]					76.6 [†]	77.2 [†]
Num shots	25	0	0	0	0	25	5	0	0	0	0	5	5
Curated shots	No					No	Yes					No	Yes
Formulation	RC	RC	RC?	RC	RC	MC	MC/RC	RC	RC	RC	RC	MC	MC/RC
Normalization	char	char	?	char?	pmi	none	none/pmi	pmi	pmi?	pmi	pmi?	none	none/pmi

Ref	Reference citation	Ref	Reference citation
Ref1	HF Open LLM Leaderboard (Beeching et al., 2023)	Ref5	OLMo paper (Groeneveld et al., 2024)
Ref2	Llama2 paper (Touvron et al., 2023a)	Ref6	Llama3 model card (AI@Meta, 2024)
Ref3	Mistral 7B (Jiang et al., 2023)	Ref7	Gemma paper (Gemma Team et al., 2024)
Ref4	Falcon paper (Almazrouei et al., 2023)	Ref8	HELM Lite Leaderboard (Liang et al., 2023)

BetterBench – Reuel et al.

- Methodology for evaluating quality of benchmarks
- “14 out of 24 benchmarks did not perform multiple evaluations of the same model or report statistical significance or uncertainty of results”
- “an average score of 5.62 [out of 15] on *Reporting statistical significance*”





Figure 7: Design and usability score for all 24 assessed benchmarks. The usability score is the weighted average of the implementation, documentation, and maintenance scores. Benchmarks were split into foundation model and non-foundation model benchmarks, depending on the model group they're targeting.

What to report in your evals


- Accuracy (overall and per-task) averaged over n runs



See Neurips slides by
Hermann et al. (2024) for
more

What to report in your evals


- Accuracy (overall and per-task) averaged over n runs
- With error bars:
 - Using standard error rather than standard deviation:
 - Uncertainty in the estimate of the mean, not just the spread of data



See Neurips slides by
Hermann et al. (2024) for
more

What to report in your evals


- Accuracy (overall and per-task) averaged over n runs
- With error bars:
 - Using standard error rather than standard deviation:
 - Uncertainty in the estimate of the mean, not just the spread of data
- How can we represent the facts that
 - Each benchmark/task might be of different quality?
 - Each model may be more/less capable/consistent on a different subset of benchmarks/tasks?



See Neurips slides by
Hermann et al. (2024) for
more

What to report in your evals

- Accuracy (overall and per-task) averaged over n runs
- With error bars:
 - Using standard error rather than standard deviation:
 - Uncertainty in the estimate of the mean, not just the spread of data
- How can we represent the facts that
 - Each benchmark/task might be of different quality?
 - Each model may be more/less capable/consistent on a different subset of benchmarks/tasks?
- How can we use these error bars to consider the probability with which we can say “Model A > Model B”?



Set up

Suppose we have n questions, and our model achieves scores $\{s_i\}_{i=1}^n$ (between 0 and 1) where s_i can be decomposed into a mean component x_i and a zero-mean random component ϵ_i .

$$s_i = x_i + \epsilon_i$$

We want to infer $\mu = \mathbb{E}[s] = \mathbb{E}[x]$

Frequentist Error Bars

Let $\bar{s} = \frac{1}{n} \sum_i s_i$. Then by the law of large numbers: $\hat{\mu} = \bar{s}$

Frequentist Error Bars

Let $\bar{s} = \frac{1}{n} \sum_i s_i$. Then by the law of large numbers: $\hat{\mu} = \bar{s}$

CLT: $SE_{CLT} = \sqrt{\text{Var}(s)/n} = \sqrt{\left(\frac{1}{n-1} \sum_i (s_i - \bar{s})^2 \right) / n}$

Frequentist Error Bars

Let $\bar{s} = \frac{1}{n} \sum_i s_i$. Then by the law of large numbers: $\hat{\mu} = \bar{s}$

CLT: $SE_{CLT} = \sqrt{\text{Var}(s)/n} = \sqrt{\left(\frac{1}{n-1} \sum_i (s_i - \bar{s})^2\right) / n}$

If $s_i \in \{0, 1\}$ then this becomes $SE_{\text{Bernoulli}} = \sqrt{\hat{s}(1 - \hat{s})/n}$

Frequentist Error Bars

Let $\bar{s} = \frac{1}{n} \sum_i s_i$. Then by the law of large numbers: $\hat{\mu} = \bar{s}$

CLT: $SE_{CLT} = \sqrt{\text{Var}(s)/n} = \sqrt{\left(\frac{1}{n-1} \sum_i (s_i - \bar{s})^2\right) / n}$

If $s_i \in \{0, 1\}$ then this becomes $SE_{\text{Bernoulli}} = \sqrt{\hat{s}(1 - \hat{s})/n}$

Obtain 95% confidence intervals as $CI_{95\%} = \bar{s} \pm 1.96 \times SE_{CLT}$

(if you do the data collection 100 times, 95 of those times you'll get $\mu \in CI_{95\%}$)

Anthropic Paper “Adding Error Bars to Evals” (Miller, 2024)

Points out that the LLama 3 (Dubey et al., 2024) paper reports only $SE_{\text{Bernoulli}}$, even when σ can take on values in $(0, 1)$ (e.g. F1 score). This leads to confidence intervals that are too wide.

 Meta

The Llama 3 Herd of Models

Llama Team, AI @ Meta¹

Anthropic Paper “Adding Error Bars to Evals” (Miller, 2024)

Points out that the Llama 3 (Dubey et al., 2024) paper reports only $SE_{\text{Bernoulli}}$, even when σ can take on values in (e.g.) F1 score). This leads to confidence intervals that are too wide.

 Meta

The Llama 3 Herd of Models

Llama Team, AI @ Meta¹

Significance estimates. Benchmark scores are estimates of a model’s true performance. These estimates have variance because benchmark sets are finite samples drawn from some underlying distribution. We follow [Madaan et al. \(2024b\)](#) and report on this variance via 95% confidence intervals (CIs), assuming that benchmark scores are Gaussian distributed. While this assumption is incorrect (*e.g.*, benchmark scores are bounded), preliminary bootstrap experiments suggest CIs (for discrete metrics) are a good approximation:

$$CI(S) = 1.96 \times \sqrt{\frac{S \times (1 - S)}{N}}.$$

Herein, S is the observed benchmark score (*e.g.*, accuracy or EM) and N the sample size of the benchmark. We omit CIs for benchmark scores that are not simple averages. We note that because subsampling is not the only source of variation, our CI values lower bound the actual variation in the capability estimate.

Anthropic Paper “Adding Error Bars to Evals” (Miller, 2024)

In reality, questions come in groups from different tasks (e.g. `boolean_expression`, `date_understanding` from BBH), so the IID assumption of the CLT is violated.

Anthropic Paper “Adding Error Bars to Evals” (Miller, 2024)

In reality, questions come in groups from different tasks (e.g. `boolean_expression`, `date_understanding` from BBH), so the IID assumption of the CLT is violated.

Therefore suggests using clustered standard error, where $s_{i,c}$ is the i th question from cluster/task c

$$\text{SE}_{\text{clustered}} = \left(\text{SE}_{\text{C.L.T.}}^2 + \frac{1}{n^2} \sum_c \sum_i \sum_{j \neq i} (s_{i,c} - \bar{s})(s_{j,c} - \bar{s}) \right)^{1/2}$$

Anthropic Paper “Adding Error Bars to Evals” (Miller, 2024)

In reality, questions come in groups from different tasks (e.g. `boolean_expression`, `date_understanding` from BBH), so the IID assumption of the CLT is violated.

Therefore suggests using clustered standard error, where $s_{i,c}$ is the i th question from cluster/task c

$$\text{SE}_{\text{clustered}} = \left(\text{SE}_{\text{C.L.T.}}^2 + \frac{1}{n^2} \sum_c \sum_i \sum_{j \neq i} (s_{i,c} - \bar{s})(s_{j,c} - \bar{s}) \right)^{1/2}$$

	$\text{SE}_{\text{clustered}}$	$\text{SE}_{\text{C.L.T.}}$	Ratio
DROP	(1.34)	(0.44)	3.05
RACE-H	(0.51%)	(0.46%)	1.10
MGSM	(1.62%)	(0.86%)	1.88

Table 4: Clustered and naive standard errors computed on two popular evals using Anthropic models (non-fictional numbers). Analyzing the same data, clustered standard errors can be over 3X larger than naive standard errors.

“sliding scale between cases where scores within a cluster are perfectly correlated ([...] each cluster acts as a single indep. observation) and perfectly uncorrelated”

Anthropic Paper “Adding Error Bars to Evals” (Miller, 2024)

How to compare two models, A and B?

Naively: $\hat{\mu}_{A-B} = \hat{\mu}_A - \hat{\mu}_B$ $SE_{A-B} = \sqrt{SE_A^2 + SE_B^2}$ $CI_{A-B,95\%} = \hat{\mu}_{A-B} \pm 1.96 \times SE_{A-B}$

Anthropic Paper “Adding Error Bars to Evals” (Miller, 2024)

How to compare two models, A and B?

Naively: $\hat{\mu}_{A-B} = \hat{\mu}_A - \hat{\mu}_B$ $SE_{A-B} = \sqrt{SE_A^2 + SE_B^2}$ $CI_{A-B,95\%} = \hat{\mu}_{A-B} \pm 1.96 \times SE_{A-B}$

But if you know the set of questions used for A and B’s evals you can do paired analysis: $s_{A-B,i} = s_{A,i} - s_{B,i}$ $\bar{s}_{A-B} = \bar{s}_A - \bar{s}_B$

$$SE_{A-B,\text{paired}} = \sqrt{\text{Var}(s_{A-B})/n} = \sqrt{\left(\frac{1}{n-1} \sum_i (s_{A-B,i} - \bar{s}_{A-B})^2 \right) / n}$$

Anthropic Paper “Adding Error Bars to Evals” (Miller, 2024)

How to compare two models, A and B?

Naively: $\hat{\mu}_{A-B} = \hat{\mu}_A - \hat{\mu}_B$ $SE_{A-B} = \sqrt{SE_A^2 + SE_B^2}$ $CI_{A-B,95\%} = \hat{\mu}_{A-B} \pm 1.96 \times SE_{A-B}$

But if you know the set of questions used for A and B’s evals you can do paired analysis: $s_{A-B,i} = s_{A,i} - s_{B,i}$ $\bar{s}_{A-B} = \bar{s}_A - \bar{s}_B$

$$SE_{A-B,\text{paired}} = \sqrt{\text{Var}(s_{A-B})/n} = \sqrt{\left(\frac{1}{n-1} \sum_i (s_{A-B,i} - \bar{s}_{A-B})^2 \right) / n}$$

Or with clustering:

$$SE_{A-B,\text{paired,clustered}} = \frac{1}{n} \left(\sum_c \sum_i \sum_j (s_{A-B,i,c} - \bar{s}_{A-B})(s_{A-B,j,c} - \bar{s}_{A-B}) \right)^{1/2}$$

How are we modelling the process of evals?

So far just used CLT/Gaussianity assumption – can we do better?

Anthropic paper suggests (briefly) $x_i \sim \mathcal{U}[0, 1]$ and $s_i \sim \text{Ber}(x_i)$ so $\epsilon_i = 1 - x_i$ with probability x_i and $\epsilon_i = -x_i$ with probability $1 - x_i$.

How are we modelling the process of evals?

So far just used CLT/Gaussianity assumption – can we do better?

Anthropic paper suggests (briefly) $x_i \sim \mathcal{U}[0, 1]$ and $s_i \sim \text{Ber}(x_i)$ so $\epsilon_i = 1 - x_i$ with probability x_i and $\epsilon_i = -x_i$ with probability $1 - x_i$.

Desi's blogs go into more detail with Binomial modelling.

Desi's Blog

Examines “GSM-Symbolic: Understanding the Limitations of Mathematical Reasoning in Large Language Models” by Mirzadeh et al. (2024) which:

- Takes the popular maths benchmark GSM8K and creates a new version GSM-Symbolic in which we can have “same questions, different numbers”

GSM8K

When Sophie watches her nephew, she gets out a variety of toys for him. The bag of building blocks has 31 blocks in it. The bin of stuffed animals has 8 stuffed animals inside. The tower of stacking rings has 9 multicolored rings on it. Sophie recently bought a tube of bouncy balls, bringing her total number of toys for her nephew up to 62. How many bouncy balls came in the tube?

Let T be the number of bouncy balls in the tube.
After buying the tube of balls, Sophie has $31+8+9+T = 48+T=62$ toys for her nephew.
Thus, $T = 62-48 = <<62-48-14>>14$ bouncy balls came in the tube.

GSM Symbolic Template

When {name} watches her {family}, she gets out a variety of toys for him. The bag of building blocks has {x} blocks in it. The bin of stuffed animals has {y} stuffed animals inside. The tower of stacking rings has {z} multicolored rings on it. {name} recently bought a tube of bouncy balls, bringing her total number of toys she bought for her {family} up to {total}. How many bouncy balls came in the tube?

```
#variables:  
- name = sample(names)  
- family = sample(["nephew", "cousin", "brother"])  
- x = range(5, 100)  
- y = range(5, 100)  
- z = range(5, 100)  
- total = range(100, 500)  
- ans = range(85, 200)
```

```
#conditions:  
- x + y + z + ans == total
```

Let T be the number of bouncy balls in the tube. After buying the tube of balls, {name} has {x} + {y} + {z} + T = {x + y + z} + T = {total} toys for her {family}.

Thus, $T = \{total\} - \{x + y + z\} = <<\{total\} - \{x + y + z\}\> \{ans\}$ bouncy balls came in the tube.

Desi's Blog

Examines “GSM-Symbolic: Understanding the Limitations of Mathematical Reasoning in Large Language Models” by Mirzadeh et al. (2024) which:

- Takes the popular maths benchmark GSM8K and creates a new version GSM-Symbolic in which we can have “same questions, different numbers”
- Claims that:
 1. Models exhibit worse performance on GSM-Symbolic than GSM8K; and
 2. This implies that “current LLMs are not capable of genuine logical reasoning; instead, they attempt to replicate the reasoning steps observed in their training data”

GSM8K

When Sophie watches her nephew, she gets out a variety of toys for him. The bag of building blocks has 31 blocks in it. The bin of stuffed animals has 8 stuffed animals inside. The tower of stacking rings has 9 multicolored rings on it. Sophie recently bought a tube of bouncy balls, bringing her total number of toys for her nephew up to 62. How many bouncy balls came in the tube?

Let T be the number of bouncy balls in the tube.
After buying the tube of balls, Sophie has $31+8+9+T = 48+T=62$ toys for her nephew.
Thus, $T = 62-48 = <<62-48-14>>14$ bouncy balls came in the tube.

GSM Symbolic Template

When {name} watches her {family}, she gets out a variety of toys for him. The bag of building blocks has {x} blocks in it. The bin of stuffed animals has {y} stuffed animals inside. The tower of stacking rings has {z} multicolored rings on it. {name} recently bought a tube of bouncy balls, bringing her total number of toys she bought for her {family} up to {total}. How many bouncy balls came in the tube?

```
#variables:  
- name = sample(names)  
- family = sample(["nephew", "cousin", "brother"])  
- x = range(5, 100)  
- y = range(5, 100)  
- z = range(5, 100)  
- total = range(100, 500)  
- ans = range(85, 200)
```

```
#conditions:  
- x + y + z + ans == total
```

Let T be the number of bouncy balls in the tube. After buying the tube of balls, {name} has {x} + {y} + {z} + T = {x + y + z} + T = {total} toys for her {family}.

Thus, $T = \{total\} - \{x + y + z\} = <<\{total\}-\{x + y + z\}\>=\{ans\}>>\{ans\}$ bouncy balls came in the tube.

But claim 1 needs backing up

- Mirzadeh et al. create 50 versions of GSM-Symbolic and run evals for various models
- Compare against a point estimate for GSM8K
- Assumption: if a model is doing 'logic' then its performance on these datasets should be equivalent.

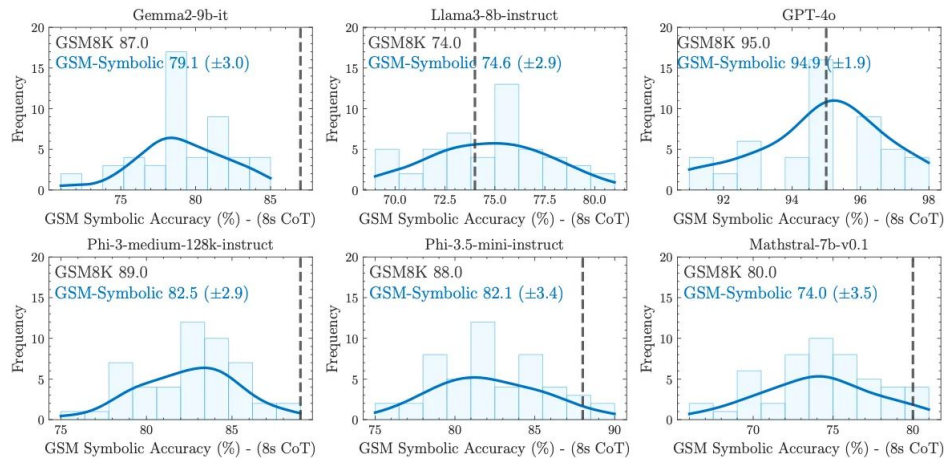


Figure 2: The distribution of 8-shot Chain-of-Thought (CoT) performance across 50 sets generated from GSM-Symbolic templates shows significant variability in accuracy among all state-of-the-art models. Furthermore, for most models, the average performance on GSM-Symbolic is lower than on GSM8K (indicated by the dashed line). Interestingly, the performance of GSM8K falls on the right side of the distribution, which, statistically speaking, should have a very low likelihood, given that GSM8K is basically a single draw from GSM-Symbolic.

But claim 1 needs backing up


Assume each model m has a probability of success p_m (the same across all N questions).

Then the number of correct answers is $\text{Bin}(N, p_m)$.

So we can get (Wilson) confidence intervals.

Now it's less obvious that model performance on GSM-Symbolic *is* actually worse than on GSM8K.

$$\begin{aligned} n_s &= \text{\#success} \\ n_f &= \text{\#fail} \\ z_\alpha &\approx \frac{(p - \hat{p})}{\sigma_n} \\ p &\in_{\approx \alpha} \frac{n_s + \frac{1}{2} z_\alpha^2}{n + z_\alpha^2} \pm \frac{z_\alpha}{n + z_\alpha^2} \sqrt{\frac{n_s n_f}{n} + \frac{z_\alpha^2}{4}} \end{aligned}$$



Is the drop in performance statistically significant?

Let p_m^{8K} be the true probability of success for model m on GSM8K and p_m^{Symb} be the true probability of success on GSM-Symbolic.

Is the drop in performance statistically significant?

Let p_m^{8K} be the true probability of success for model m on GSM8K and p_m^{Symb} be the true probability of success on GSM-Symbolic.

Run Fisher's exact test on two-sided and one-sided hypotheses:

$$H_0 : p_m^{8K} = p_m^{\text{Symb}}$$

$$H_1^{\text{two-sided}} : p_m^{8K} \neq p_m^{\text{Symb}}$$


$$H_0 : p_m^{8K} = p_m^{\text{Symb}}$$

$$H_1^{\text{one-sided}} : p_m^{8K} > p_m^{\text{Symb}}$$

Is the drop in performance statistically significant? ... Maybe

It turns out there is some (weak) evidence that the models *are* performing worse on GSM-Symbolic compared to GSM8K

(especially when you include irrelevant information in the questions)



Bayesian Error Bars

- Can we do similar eval analysis/modelling with Bayes?
 - Yes!

Bayesian Error Bars

- Can we do similar eval analysis/modelling with Bayes?
 - Yes!
- Benefits:
 - Rather than confidence intervals (95% of which may contain the true parameter) we can use credible intervals and make direct statements such as “the true parameter lies in here with 95% probability”.
 - Allow for complex hierarchical modelling
 - Doesn't rely on CLT (IID and large N assumptions)
 - Avoid complicated construction of frequentist setups (p-values shenanigans)

Bayesian Error Bars

- Can we do similar eval analysis/modelling with Bayes?
 - Yes!
- Benefits:
 - Rather than confidence intervals (95% of which may contain the true parameter) we can use credible intervals and make direct statements such as “the true parameter lies in here with 95% probability”.
 - Allow for complex hierarchical modelling
 - Doesn't rely on CLT (IID and large N assumptions)
 - Avoid complicated construction of frequentist setups (p-values shenanigans)
- Drawbacks:
 - Have to choose priors carefully (actually not too much of an issue)
 - Computation (e.g. Stan)

Experiments Setup

- 29 LLMs
 - actually have ~60
- 24 Tasks (BigBench Hard) (about 200-250 questions each)
 - actually have at least ~40 tasks
- Raw eval binary data (success/fail) from Huggingface leaderboard \mathcal{D}



```
meta-llama/Meta-Llama-3.1-70B
meta-llama/Meta-Llama-3.1-8B
meta-llama/Meta-Llama-3-8B
meta-llama/Meta-Llama-3-70B
mistralai/Mixtral-8x7B-v0.1
mistralai/Mixtral-8x22B-v0.1
google/gemma-2-27b
google/gemma-2-2b
Qwen/Qwen1.5-7B
Qwen/Qwen1.5-110B
Qwen/Qwen2.5-0.5B
Qwen/Qwen2.5-3B
Qwen/Qwen2.5-14B
Qwen/Qwen2.5-72B
meta-llama/Meta-Llama-3.1-70B-Instruct
meta-llama/Meta-Llama-3.1-8B-Instruct
meta-llama/Meta-Llama-3-8B-Instruct
meta-llama/Meta-Llama-3-70B-Instruct
microsoft/Phi-3.5-mini-instruct
microsoft/Phi-3.5-MoE-instruct
mistralai/Mixtral-8x7B-Instruct-v0.1
mistralai/Mixtral-7B-Instruct-v0.3
google/gemma-2-27b-it
Qwen/Qwen1.5-7B-Chat
Qwen/Qwen1.5-110B-Chat
Qwen/Qwen2.5-0.5B-Instruct
Qwen/Qwen2.5-3B-Instruct
Qwen/Qwen2.5-14B-Instruct
Qwen/Qwen2.5-72B-Instruct
```


```
"bbh_boolean_expressions",
"bbh_causal_judgement",
"bbh_date_understanding",
"bbh_disambiguation_qa",
"bbh_formal_fallacies",
"bbh_geometric_shapes",
"bbh_hyperbaton",
"bbh_logical_deduction_five_objects",
"bbh_logical_deduction_seven_objects",
"bbh_logical_deduction_three_objects",
"bbh_movie_recommendation",
"bbh_navigate",
"bbh_object_counting",
"bbh_penguins_in_a_table",
"bbh_reasoning_about_colored_objects",
"bbh_ruin_names",
"bbh_salient_translation_error_detection",
"bbh_snarks",
"bbh_sports_understanding",
"bbh_temporal_sequences",
"bbh_tracking_shuffled_objects_five_objects",
"bbh_tracking_shuffled_objects_seven_objects",
"bbh_tracking_shuffled_objects_three_objects",
"bbh_web_of_lies",
```


Model 1: Beta Binomial

Prior $p_m \sim \text{Beta}(1, 1)$

Model $\bar{s}_{m,t} \sim \text{Bin}(N_t, p_m)$

Infer $\{p_m\}_{m=1}^M | \mathcal{D}$

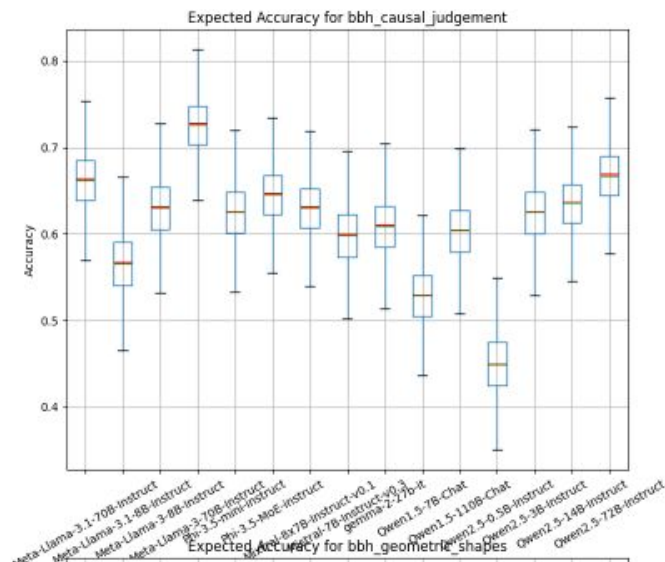
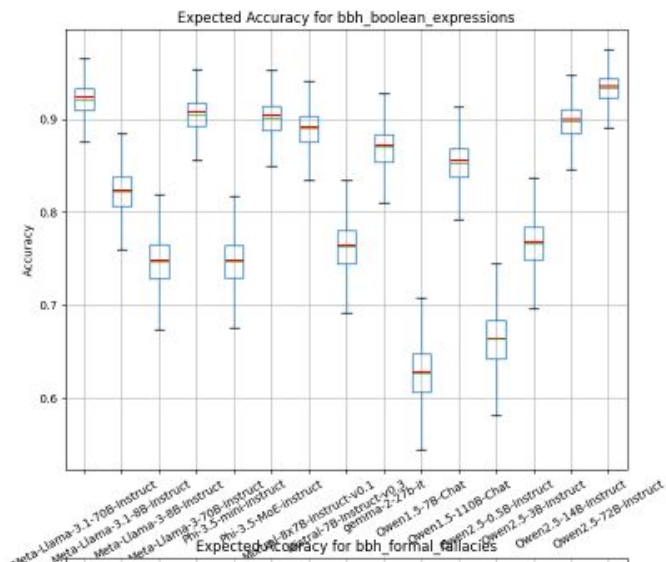


Model 2: Beta Binomial with per-task model performance

Prior $p_{m,t} \sim \text{Beta}(1, 1)$

Model $\bar{s}_{m,t} \sim \text{Bin}(N_t, p_{m,t})$

Infer $\{p_{m,t}\}_{m=1}^M | \mathcal{D}$



Model 3: Question-based model

$$\text{bias} \sim \mathcal{N}(0, 1)$$

$$\text{taskDifficultyStd} \sim \text{Gamma}(2, 2)$$

$$\text{questionDifficultyStd} \sim \text{Gamma}(2, 2)$$

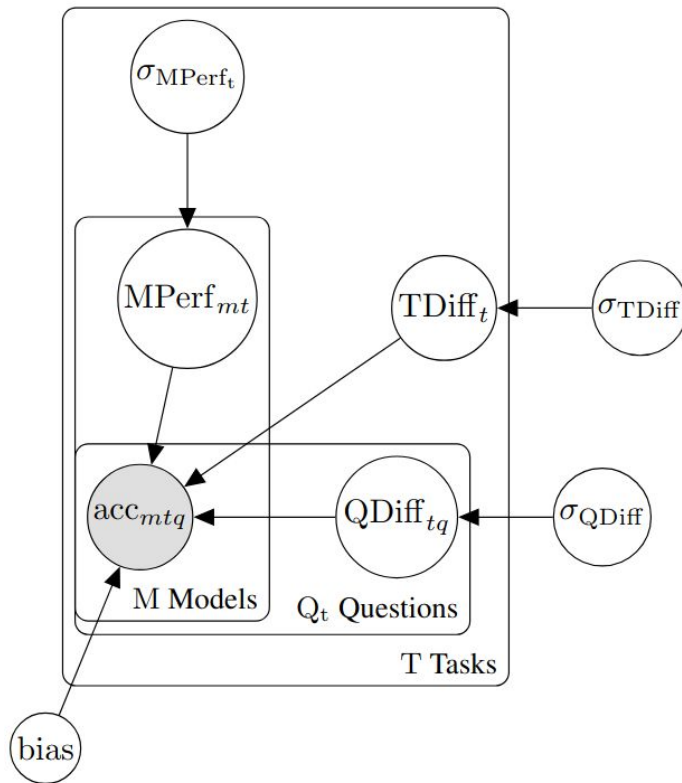
$$\text{modelPerformanceStd} \sim \text{Gamma}(2, 2)$$

$$\text{taskDifficulty}_t \sim \mathcal{N}(0, \text{taskDifficultyStd})$$

$$\text{questionDifficulty}_q \sim \mathcal{N}(0, \text{questionDifficultyStd})$$

$$\text{modelPerformance}_t \sim \mathcal{N}(0, \text{modelPerformanceStd})$$

$$\begin{aligned} \bar{s}_{m,q} \sim \text{Ber}(\text{sigmoid}(\text{bias} & \\ & + \text{modelPerformance}_m \\ & - \text{taskDifficulty}_{\text{task}(q)} \\ & - \text{questionDifficulty}_q)) \end{aligned}$$



Model 3: Question-based model

$$\text{bias} \sim \mathcal{N}(0, 1)$$

$$\text{taskDifficultyStd} \sim \text{Gamma}(2, 2)$$

$$\text{questionDifficultyStd} \sim \text{Gamma}(2, 2)$$



$$\text{modelPerformanceStd} \sim \text{Gamma}(2, 2)$$

$$\text{taskDifficulty}_t \sim \mathcal{N}(0, \text{taskDifficultyStd})$$

$$\text{questionDifficulty}_q \sim \mathcal{N}(0, \text{questionDifficultyStd})$$

$$\text{modelPerformance}_t \sim \mathcal{N}(0, \text{modelPerformanceStd})$$

$$\begin{aligned} \bar{s}_{m,q} \sim & \text{Ber}(\text{sigmoid}(\text{bias} \\ & + \text{modelPerformance}_m \\ & - \text{taskDifficulty}_{\text{task}(q)} \\ & - \text{questionDifficulty}_q)) \end{aligned}$$



Model 4: Question-based model w/ across-task performance

$$\text{bias} \sim \mathcal{N}(0, 1)$$

$$\text{taskDifficultyStd} \sim \text{Gamma}(2, 2)$$

$$\text{questionDifficultyStd} \sim \text{Gamma}(2, 2)$$

$$\text{modelPerformanceStd} \sim \text{Gamma}(2, 2)$$

$$\text{acrossTaskPerformanceStd} \sim \text{Gamma}(2, 2)$$

$$\text{taskDifficulty}_t \sim \mathcal{N}(0, \text{taskDifficultyStd})$$

$$\text{questionDifficulty}_q \sim \mathcal{N}(0, \text{questionDifficultyStd})$$

$$\text{modelPerformance}_{m,t} \sim \mathcal{N}(0, \text{modelPerformanceStd})$$

$$\text{acrossTaskPerformance}_m \sim \mathcal{N}(0, \text{acrossTaskPerformanceStd})$$

$$\begin{aligned} \bar{s}_{m,q} \sim & \text{Ber}(\text{sigmoid}(\text{bias} \\ & + \text{modelPerformance}_{m,t} \\ & + \text{acrossTaskPerformance}_m \\ & - \text{taskDifficulty}_{\text{task}(q)} \\ & - \text{questionDifficulty}_q)) \end{aligned}$$

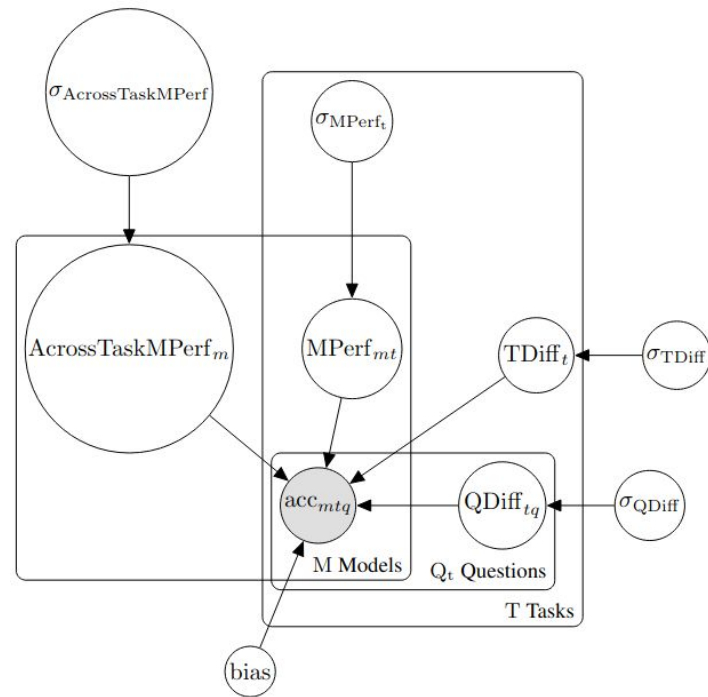


Figure 4. Graphical model for the hierarchical model with per-model overall performance variables.

Model 3: w/ across-task performance

bias $\sim \mathcal{N}(0, 1)$

taskDifficultyStd $\sim \text{Gamma}(2, 2)$

questionDifficultyStd $\sim \text{Gamma}(2, 2)$

modelPerformanceStd $\sim \text{Gamma}(2, 2)$

acrossTaskPerformanceStd $\sim \text{Gamma}(2, 2)$

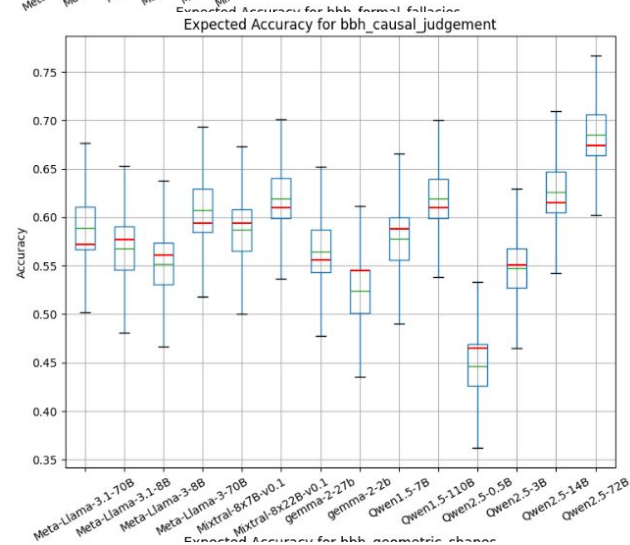

taskDifficulty_t $\sim \mathcal{N}(0, \text{taskDifficultyStd})$

questionDifficulty_q $\sim \mathcal{N}(0, \text{questionDifficultyStd})$

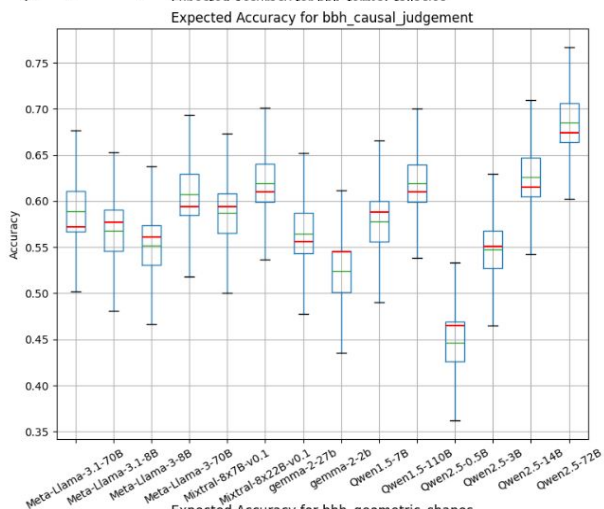

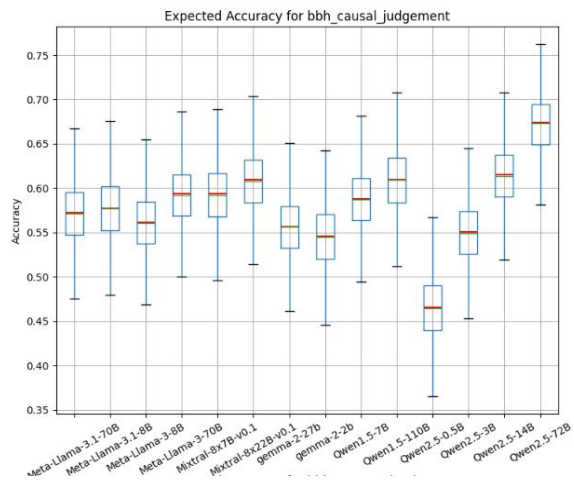
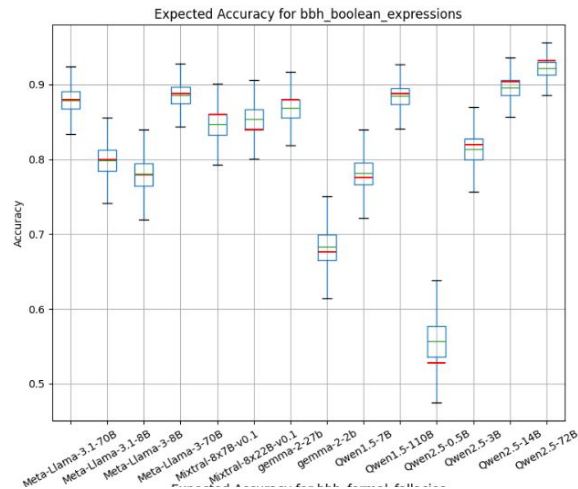
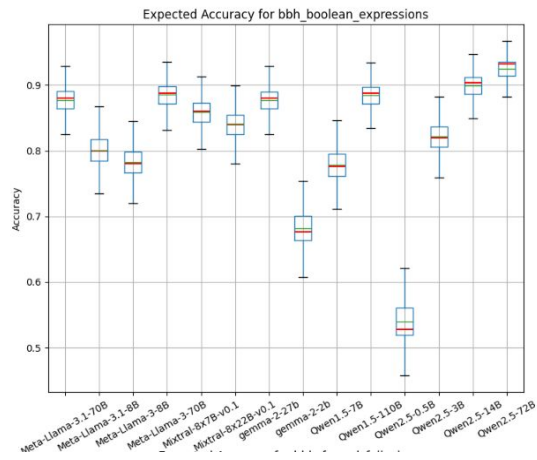

modelPerformance_{m,t} $\sim \mathcal{N}(0, \text{modelPerformanceStd})$

acrossTaskPerformance_m $\sim \mathcal{N}(0, \text{acrossTaskPerformanceStd})$

$$\bar{s}_{m,q} \sim \text{Ber}(\text{sigmoid}(\text{bias} + \text{modelPerformance}_{m,t} + \text{acrossTaskPerformance}_m - \text{taskDifficulty}_{\text{task}(q)} - \text{questionDifficulty}_q))$$



Model 2 vs 3 vs 4 - some (slight) reduction in size of error bars – needs frequentist comparison



Bayesian hierarchical modelling for evals makes it easy to...

1. Easily find a $\text{Prob}(\text{Model A} > \text{Model B})$ – HMC gives us samples of $p_{m,t}|\mathcal{D}$, just look at

$$\frac{1}{\text{num samples}} \sum_{\text{samples}} \frac{1}{T} \sum_t \mathbb{I}[p_{A,t} > p_{B,t}]$$

Bayesian hierarchical modelling for evals makes it easy to...

1. Easily find a $\text{Prob}(\text{Model A} > \text{Model B})$ – HMC gives us samples of $p_{m,t}|\mathcal{D}$, just look at

$$\frac{1}{\text{num samples}} \sum_{\text{samples}} \frac{1}{T} \sum_t \mathbb{I}[p_{A,t} > p_{B,t}]$$

2. Elicit information about latent variables of interest:
 - a. Task difficulty
 - b. Per-task model performance
 - c. Capabilities*

Bayesian hierarchical modelling for evals makes it easy to...

1. Easily find a $\text{Prob}(\text{Model A} > \text{Model B})$ – HMC gives us samples of $p_{m,t}|\mathcal{D}$, just look at

$$\frac{1}{\text{num samples}} \sum_{\text{samples}} \frac{1}{T} \sum_t \mathbb{I}[p_{A,t} > p_{B,t}]$$

2. Elicit information about latent variables of interest:
 - a. Task difficulty
 - b. Per-task model performance
 - c. Capabilities*

* If we make modelPerformance and taskDifficulty vectors of length C (say 5 or 7) with likelihood s.t. $\bar{s}_{m,t}$ depends on their dot product, each dimension of those vectors might be interpretable as model/task capabilities e.g. arithmetic, grammar, logic etc.

References

- Anthropic. 2023. 'Challenges in Evaluating AI Systems'. 4 October 2023. <https://www.anthropic.com/news/evaluating-ai-systems>.
- Biderman, Stella, Hailey Schoelkopf, Lintang Sutawika, Leo Gao, Jonathan Tow, Baber Abbasi, Alham Fikri Aji, et al. 2024. 'Lessons from the Trenches on Reproducible Evaluation of Language Models'. arXiv. <https://doi.org/10.48550/arXiv.2405.14782>.
- Grattafiori, Aaron, Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, Ahmad Al-Dahle, Aiesha Letman, et al. 2024. 'The Llama 3 Herd of Models'. arXiv. <https://doi.org/10.48550/arXiv.2407.21783>.
- Gu, Yuling, Oyvind Tafjord, Bailey Kuehl, Dany Haddad, Jesse Dodge, and Hannaneh Hajishirzi. 2024. 'OLMES: A Standard for Language Model Evaluations'. arXiv. <https://doi.org/10.48550/arXiv.2406.08446>.
- Hermann, Katherine, Jennifer Hu, and Michael Mozer. n.d. 'Experimental Design and Analysis for AI Researchers'.
- Ivanova, Desi. 2024a. 'On Some (Fixable) Limitations of "Understanding the Limitations of Mathematical Reasoning in LLMs"'. 22 October 2024. https://substack.com/inbox/post/150508215?utm_campaign=post&triedRedirect=true.
- ———. 2024b. 'Towards More Rigorous Evaluations of Language Models'. 28 November 2024. <https://substack.com/home/post/p-152149873>.
- Liang, Percy, Rishi Bommasani, Tony Lee, Dimitris Tsipras, Dilara Soylu, Michihiro Yasunaga, Yian Zhang, et al. 2023. 'Holistic Evaluation of Language Models'. arXiv. <https://doi.org/10.48550/arXiv.2211.09110>.
- Miller, Evan. 2024. 'Adding Error Bars to Evals: A Statistical Approach to Language Model Evaluations'. arXiv. <https://doi.org/10.48550/arXiv.2411.00640>.
- Mirzadeh, Iman, Keivan Alizadeh, Hooman Shahrokhi, Oncl Tuzel, Samy Bengio, and Mehrdad Farajtabar. 2024. 'GSM-Symbolic: Understanding the Limitations of Mathematical Reasoning in Large Language Models'. arXiv. <https://doi.org/10.48550/arXiv.2410.05229>.
- Perez, Ethan, Sam Ringer, Kamilė Lukošūtė, Karina Nguyen, Edwin Chen, Scott Heiner, Craig Pettit, et al. 2022. 'Discovering Language Model Behaviors with Model-Written Evaluations'. arXiv. <https://doi.org/10.48550/arXiv.2212.09251>.
- Polo, Felipe Maia, Lucas Weber, Leshem Choshen, Yuekai Sun, Gongjun Xu, and Mikhail Yurochkin. 2024. 'tinyBenchmarks: Evaluating LLMs with Fewer Examples'. arXiv. <https://doi.org/10.48550/arXiv.2402.14992>.
- Reuel, Anka, Amelia Hardy, Chandler Smith, Max Lamparth, Malcolm Hardy, and Mykel J. Kochenderfer. 2024. 'BetterBench: Assessing AI Benchmarks, Uncovering Issues, and Establishing Best Practices'. arXiv. <https://doi.org/10.48550/arXiv.2411.12990>.
- Song, Yifan, Guoyin Wang, Sujian Li, and Bill Yuchen Lin. 2024. 'The Good, The Bad, and The Greedy: Evaluation of LLMs Should Not Ignore Non-Determinism'. arXiv. <https://doi.org/10.48550/arXiv.2407.10457>.