

Q-SFT: Q-Learning as Supervised Fine-Tuning

12 August 2025

Motivation

- Value-based RL:

Motivation

- Value-based RL:
 - Compute the expected future reward of a state-action pair $Q(s, a)$, make decisions based on these values

Motivation

- Value-based RL:
 - Compute the expected future reward of a state-action pair $Q(s, a)$, make decisions based on these values
 - (Rather than learning a policy directly without a value function)

Motivation

- Value-based RL:
 - Compute the expected future reward of a state-action pair $Q(s, a)$, make decisions based on these values
 - (Rather than learning a policy directly without a value function)
 - Can be more sample-efficient than policy gradient methods

Motivation

- Value-based RL:
 - Compute the expected future reward of a state-action pair $Q(s, a)$, make decisions based on these values
 - (Rather than learning a policy directly without a value function)
 - Can be more sample-efficient than policy gradient methods
- But, applying value-based RL to language models is challenging:

Motivation

- Value-based RL:
 - Compute the expected future reward of a state-action pair $Q(s, a)$, make decisions based on these values
 - (Rather than learning a policy directly without a value function)
 - Can be more sample-efficient than policy gradient methods
- But, applying value-based RL to language models is challenging:
 - The action and state spaces are very large

Motivation

- Value-based RL:
 - Compute the expected future reward of a state-action pair $Q(s, a)$, make decisions based on these values
 - (Rather than learning a policy directly without a value function)
 - Can be more sample-efficient than policy gradient methods
- But, applying value-based RL to language models is challenging:
 - The action and state spaces are very large
 - The reward signal is sparse and delayed

Motivation

- Value-based RL:
 - Compute the expected future reward of a state-action pair $Q(s, a)$, make decisions based on these values
 - (Rather than learning a policy directly without a value function)
 - Can be more sample-efficient than policy gradient methods
- But, applying value-based RL to language models is challenging:
 - The action and state spaces are very large
 - The reward signal is sparse and delayed
 - Training is often unstable

Motivation

- Value-based RL:
 - Compute the expected future reward of a state-action pair $Q(s, a)$, make decisions based on these values
 - (Rather than learning a policy directly without a value function)
 - Can be more sample-efficient than policy gradient methods
- But, applying value-based RL to language models is challenging:
 - The action and state spaces are very large
 - The reward signal is sparse and delayed
 - Training is often unstable
- The idea: use the pretrained model's output probs/logits to compute Q-values ($Q^\pi(s, a) = \text{expected future reward from action } a \text{ in state } s \text{ following the policy } \pi$).

Motivation

- Value-based RL:
 - Compute the expected future reward of a state-action pair $Q(s, a)$, make decisions based on these values
 - (Rather than learning a policy directly without a value function)
 - Can be more sample-efficient than policy gradient methods
- But, applying value-based RL to language models is challenging:
 - The action and state spaces are very large
 - The reward signal is sparse and delayed
 - Training is often unstable
- The idea: use the pretrained model's output probs/logits to compute Q-values ($Q^\pi(s, a) = \text{expected future reward from action } a \text{ in state } s \text{ following the policy } \pi$).
 - Logits incorporate high-quality information from pretraining

Motivation

- Value-based RL:
 - Compute the expected future reward of a state-action pair $Q(s, a)$, make decisions based on these values
 - (Rather than learning a policy directly without a value function)
 - Can be more sample-efficient than policy gradient methods
- But, applying value-based RL to language models is challenging:
 - The action and state spaces are very large
 - The reward signal is sparse and delayed
 - Training is often unstable
- The idea: use the pretrained model's output probs/logits to compute Q-values ($Q^\pi(s, a) =$ expected future reward from action a in state s following the policy π).
 - Logits incorporate high-quality information from pretraining
 - Doesn't require a change in architecture or a separate value function network/head

Motivation

- Value-based RL:
 - Compute the expected future reward of a state-action pair $Q(s, a)$, make decisions based on these values
 - (Rather than learning a policy directly without a value function)
 - Can be more sample-efficient than policy gradient methods
- But, applying value-based RL to language models is challenging:
 - The action and state spaces are very large
 - The reward signal is sparse and delayed
 - Training is often unstable
- The idea: use the pretrained model's output probs/logits to compute Q-values ($Q^\pi(s, a) = \text{expected future reward from action } a \text{ in state } s \text{ following the policy } \pi$).
 - Logits incorporate high-quality information from pretraining
 - Doesn't require a change in architecture or a separate value function network/head
 - Should provide a stable initialisation and training signal

Key RL Components

- **State (s):** The current dialogue history/text sequence.

Key RL Components

- **State (s)**: The current dialogue history/text sequence.
- **Action (a)**: The next token to be generated.

Key RL Components

- **State (s)**: The current dialogue history/text sequence.
- **Action (a)**: The next token to be generated.
- **Reward (r)**: Usually sparse and delayed.

Key RL Components

- **State (s)**: The current dialogue history/text sequence.
- **Action (a)**: The next token to be generated.
- **Reward (r)**: Usually sparse and delayed.
- **Discount factor ($\gamma \in [0, 1]$)**: The discount factor for future rewards.

Key RL Components

- **State (s)**: The current dialogue history/text sequence.
- **Action (a)**: The next token to be generated.
- **Reward (r)**: Usually sparse and delayed.
- **Discount factor ($\gamma \in [0, 1]$)**: The discount factor for future rewards.
 - (Prioritizes immediate rewards over delayed ones.)

Key RL Components

- **State (s)**: The current dialogue history/text sequence.
- **Action (a)**: The next token to be generated.
- **Reward (r)**: Usually sparse and delayed.
- **Discount factor ($\gamma \in [0, 1]$)**: The discount factor for future rewards.
 - (Prioritizes immediate rewards over delayed ones.)
 - (We can set it to 1 and ignore it for single-turn interactions with rewards only on the final token.)

Key RL Components

- **State (s):** The current dialogue history/text sequence.
- **Action (a):** The next token to be generated.
- **Reward (r):** Usually sparse and delayed.
- **Discount factor ($\gamma \in [0, 1]$):** The discount factor for future rewards.
 - (Prioritizes immediate rewards over delayed ones.)
 - (We can set it to 1 and ignore it for single-turn interactions with rewards only on the final token.)
- **Policy (π):** The policy we are trying to improve (distribution over next-tokens).

The Bellman Equation and Q-Learning

- **Q-function (Q^π):** The discounted expected future reward from a state-action pair following policy π .

The Bellman Equation and Q-Learning

- **Q-function (Q^π):** The discounted expected future reward from a state-action pair following policy π .
- Q^π satisfies the Bellman recurrence:

$$Q^\pi(s, a) = r(s, a) + \gamma \mathbb{E}_{s' \sim P(\cdot|s, a), a' \sim \pi(\cdot|s')} [Q^\pi(s', a')]$$

The Bellman Equation and Q-Learning

- **Q-function (Q^π):** The discounted expected future reward from a state-action pair following policy π .
- Q^π satisfies the Bellman recurrence:

$$Q^\pi(s, a) = r(s, a) + \gamma \mathbb{E}_{s' \sim P(\cdot|s, a), a' \sim \pi(\cdot|s')} [Q^\pi(s', a')]$$

- The Q-function of the optimal policy, Q^* , maximizes the expected future reward:

$$Q^*(s, a) = r(s, a) + \gamma \mathbb{E}_{s' \sim P(\cdot|s, a)} \left[\max_{a'} Q^*(s', a') \right]$$

The Bellman Equation and Q-Learning

- **Q-function (Q^π):** The discounted expected future reward from a state-action pair following policy π .
- Q^π satisfies the Bellman recurrence:

$$Q^\pi(s, a) = r(s, a) + \gamma \mathbb{E}_{s' \sim P(\cdot|s, a), a' \sim \pi(\cdot|s')} [Q^\pi(s', a')]$$

- The Q-function of the optimal policy, Q^* , maximizes the expected future reward:

$$Q^*(s, a) = r(s, a) + \gamma \mathbb{E}_{s' \sim P(\cdot|s, a)} \left[\max_{a'} Q^*(s', a') \right]$$

- The goal of RL is to find the optimal policy π^* :

$$\pi^* = \arg \max_{\pi} \mathbb{E}_{s \sim P(\cdot)} \left[\max_a Q^\pi(s, a) \right]$$

The Bellman Equation and Q-Learning

- **Q-function (Q^π):** The discounted expected future reward from a state-action pair following policy π .
- Q^π satisfies the Bellman recurrence:

$$Q^\pi(s, a) = r(s, a) + \gamma \mathbb{E}_{s' \sim P(\cdot|s, a), a' \sim \pi(\cdot|s')} [Q^\pi(s', a')]$$

- The Q-function of the optimal policy, Q^* , maximizes the expected future reward:

$$Q^*(s, a) = r(s, a) + \gamma \mathbb{E}_{s' \sim P(\cdot|s, a)} \left[\max_{a'} Q^*(s', a') \right]$$

- The goal of RL is to find the optimal policy π^* :

$$\pi^* = \arg \max_{\pi} \mathbb{E}_{s \sim P(\cdot)} \left[\max_a Q^\pi(s, a) \right]$$

- **Summary:** if we have Q^* , we can find π^* by taking the **argmax** over the **Q-values**.

Q-SFT: Q-LEARNING FOR LANGUAGE MODELS VIA SUPERVISED FINE-TUNING

Joey Hong

University of California, Berkeley
`{jxihong, anca, svlevine}@berkeley.edu`

Anca Dragan

Sergey Levine

The Core Idea: Q-Learning as SFT

- Reframe Q-learning as a supervised fine-tuning problem using an offline dataset $\mathcal{D} = \{(s_i, a_i, r_i)\}_{i=1}^N$.

The Core Idea: Q-Learning as SFT

- Reframe Q-learning as a supervised fine-tuning problem using an offline dataset $\mathcal{D} = \{(s_i, a_i, r_i)\}_{i=1}^N$.
- Finetune a base model to output probabilities $\hat{p}_\theta(a|s)$ that are estimates of the optimal Q-values, $Q^*(s, a)$.

The Core Idea: Q-Learning as SFT

- Reframe Q-learning as a supervised fine-tuning problem using an offline dataset $\mathcal{D} = \{(s_i, a_i, r_i)\}_{i=1}^N$.
- Finetune a base model to output probabilities $\hat{p}_\theta(a|s)$ that are estimates of the optimal Q-values, $Q^*(s, a)$.
- (At inference time, we'll have to modify \hat{p}_θ slightly to get a good (non-greedy) policy π_θ .)

A Note on the Dataset

- In text, the dataset $\mathcal{D} = \{(s_i, a_i, r_i)\}_{i=1}^N$ has actions that are tokens and states that are token sequences.

A Note on the Dataset

- In text, the dataset $\mathcal{D} = \{(s_i, a_i, r_i)\}_{i=1}^N$ has actions that are tokens and states that are token sequences.
- E.g. s_i might be the question and partial response, "Q: If a right-angled triangle has side lengths 3 and 4, what is the length of the hypotenuse? A: We can use".

A Note on the Dataset

- In text, the dataset $\mathcal{D} = \{(s_i, a_i, r_i)\}_{i=1}^N$ has actions that are tokens and states that are token sequences.
- E.g. s_i might be the question and partial response, "Q: If a right-angled triangle has side lengths 3 and 4, what is the length of the hypotenuse? A: We can use".
- And a_i might be the next token, "Pythagoras".

A Note on the Dataset

- In text, the dataset $\mathcal{D} = \{(s_i, a_i, r_i)\}_{i=1}^N$ has actions that are tokens and states that are token sequences.
- E.g. s_i might be the question and partial response, "Q: If a right-angled triangle has side lengths 3 and 4, what is the length of the hypotenuse? A: We can use".
- And a_i might be the next token, "Pythagoras".
- Only the final token can have a non-zero reward, but any token can have positive values of the Q-function ("Pythagoras" is a good action, but "Fermat" is not).

A Note on the Dataset

- In text, the dataset $\mathcal{D} = \{(s_i, a_i, r_i)\}_{i=1}^N$ has actions that are tokens and states that are token sequences.
- E.g. s_i might be the question and partial response, "Q: If a right-angled triangle has side lengths 3 and 4, what is the length of the hypotenuse? A: We can use".
- And a_i might be the next token, "Pythagoras".
- Only the final token can have a non-zero reward, but any token can have positive values of the Q-function ("Pythagoras" is a good action, but "Fermat" is not).
- We can extend this to multi-turn interactions by concatenating the previous responses to the question (in which case we may get intermediate rewards).

Behavioural Cloning

- We can assume our offline dataset $\mathcal{D} = \{(s_i, a_i, r_i)\}_{i=1}^N$ is generated by some behaviour policy π_b over tokens.

Behavioural Cloning

- We can assume our offline dataset $\mathcal{D} = \{(s_i, a_i, r_i)\}_{i=1}^N$ is generated by some behaviour policy π_b over tokens.
- We can train a policy π_ϕ to match the behaviour of the dataset by minimizing the cross-entropy loss:

$$L_{\text{CE}}(\phi) = \mathbb{E}_{(s,a) \sim \mathcal{D}} [\log \pi_\phi(a|s)]$$

Behavioural Cloning

- We can assume our offline dataset $\mathcal{D} = \{(s_i, a_i, r_i)\}_{i=1}^N$ is generated by some behaviour policy π_b over tokens.
- We can train a policy π_ϕ to match the behaviour of the dataset by minimizing the cross-entropy loss:

$$L_{\text{CE}}(\phi) = \mathbb{E}_{(s,a) \sim \mathcal{D}} [\log \pi_\phi(a|s)]$$

- This is just standard SFT, and let's us approximate $\pi_\phi \approx \pi_b$ as a reference policy later.

Weighted Cross-Entropy Loss

- We can modify the loss function to weight different tokens/actions with weights $w(s, a)$:

$$L_{WCE}(\theta) = \mathbb{E}_{(s,a) \sim \mathcal{D}} \left[w(s, a) \log p_{\theta}(a|s) + (1 - w(a, s)) \mathbb{E}_{a' \neq a} [\log p_{\theta}(a'|s)] \right]$$

Weighted Cross-Entropy Loss

- We can modify the loss function to weight different tokens/actions with weights $w(s, a)$:

$$L_{WCE}(\theta) = \mathbb{E}_{(s,a) \sim \mathcal{D}} \left[w(s, a) \log p_{\theta}(a|s) + (1 - w(a, s)) \mathbb{E}_{a' \neq a} [\log p_{\theta}(a'|s)] \right]$$

- Optimising this will give us $\hat{p}_{\theta}(s, a) \approx w(s, a) \pi_b(a|s)$.

Weighted Cross-Entropy Loss

- We can modify the loss function to weight different tokens/actions with weights $w(s, a)$:

$$L_{WCE}(\theta) = \mathbb{E}_{(s,a) \sim \mathcal{D}} \left[w(s, a) \log p_{\theta}(a|s) + (1 - w(a, s)) \mathbb{E}_{a' \neq a} [\log p_{\theta}(a'|s)] \right]$$

- Optimising this will give us $\hat{p}_{\theta}(s, a) \approx w(s, a) \pi_b(a|s)$.
- The goal is to find weights $w(s, a)$ that make $\hat{p}_{\theta}(s, a) \approx Q^*(s, a)$.

Weighted Cross-Entropy Loss

- We can modify the loss function to weight different tokens/actions with weights $w(s, a)$:

$$L_{WCE}(\theta) = \mathbb{E}_{(s,a) \sim \mathcal{D}} [w(s, a) \log p_{\theta}(a|s) + (1 - w(a, s)) \mathbb{E}_{a' \neq a} [\log p_{\theta}(a'|s)]]$$

- Optimising this will give us $\hat{p}_{\theta}(s, a) \approx w(s, a) \pi_b(a|s)$.
- The goal is to find weights $w(s, a)$ that make $\hat{p}_{\theta}(s, a) \approx Q^*(s, a)$.
- The authors use the *empirical Bellman probability operator*:

$$w(s, a) = \hat{\mathcal{B}}^* \bar{p}_{\theta}(a|s) = r + \gamma \max_{a'} \frac{\bar{p}_{\theta}(a'|s')}{\pi_b(a'|s')},$$

where r is the reward of moving from s to s' via action a , and \bar{p}_{θ} is a moving average of p_{θ} over training.

Weighted Cross-Entropy Loss

- We can modify the loss function to weight different tokens/actions with weights $w(s, a)$:

$$L_{WCE}(\theta) = \mathbb{E}_{(s,a) \sim \mathcal{D}} [w(s, a) \log p_{\theta}(a|s) + (1 - w(a, s)) \mathbb{E}_{a' \neq a} [\log p_{\theta}(a'|s)]]$$

- Optimising this will give us $\hat{p}_{\theta}(s, a) \approx w(s, a) \pi_b(a|s)$.
- The goal is to find weights $w(s, a)$ that make $\hat{p}_{\theta}(s, a) \approx Q^*(s, a)$.
- The authors use the *empirical Bellman probability operator*:

$$w(s, a) = \hat{\mathcal{B}}^* \bar{p}_{\theta}(a|s) = r + \gamma \max_{a'} \frac{\bar{p}_{\theta}(a'|s')}{\pi_b(a'|s')},$$

where r is the reward of moving from s to s' via action a , and \bar{p}_{θ} is a moving average of p_{θ} over training.

- They prove that this leads to \hat{p}_{θ} being a good (“conservative”) approximation of Q^* .

Q-SFT: Policy Extraction

We want to extract a policy $\hat{\pi}$ from the Q-SFT model.

- Greedy policy: $\hat{\pi}(a|s) = \mathbf{1}[a = \arg \max_a Q^*(s, a)]$.

Q-SFT: Policy Extraction

We want to extract a policy $\hat{\pi}$ from the Q-SFT model.

- Greedy policy: $\hat{\pi}(a|s) = \mathbf{1}[a = \arg \max_a Q^*(s, a)]$.
- Entropy-regularised policy: $\hat{\pi}(a|s) \propto \exp(Q^*(s, a))$.

Q-SFT: Policy Extraction

We want to extract a policy $\hat{\pi}$ from the Q-SFT model.

- Greedy policy: $\hat{\pi}(a|s) = \mathbf{1}[a = \arg \max_a Q^*(s, a)]$.
- Entropy-regularised policy: $\hat{\pi}(a|s) \propto \exp(Q^*(s, a))$.
- KL-regularised policy (suggested by the authors) with hyperparameter $\beta > 0$:

$$\begin{aligned}\hat{\pi}(a|s) &\propto \pi_b(a|s) \exp(\beta Q^*(s, a)) \\ &\approx \pi_\phi(a|s) \exp(\beta p_\theta(a|s))\end{aligned}$$

This is a well-known solution to the constrained optimisation problem:

$$\arg \max_{\pi} \mathbb{E}_{s \sim P(\cdot), a \sim \pi(\cdot|s)} [Q^*(s, a)] \quad s.t. \quad \mathbb{E}_{s \sim P(\cdot)} [D_{\text{KL}}(\pi(\cdot|s) \parallel \pi_b(\cdot|s))] \leq \epsilon$$

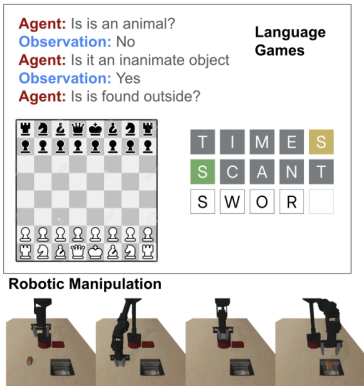
The Q-SFT Algorithm

Algorithm 1 Q-learning via Supervised Fine-tuning (Q-SFT)

Require: Dataset $\mathcal{D} = \{(s_i, a_i, r_i, s'_i)\}_{i \in [N]}$, hyperparameter $\beta > 0$

- 1: Initialize $\phi, \theta, \bar{\theta}$ from pretrained model.
 - 2: *Optimize behavior policy:*
 - 3: **for** each gradient step **do**
 - 4: Update $\phi \leftarrow \phi - \lambda_\phi \nabla_\phi \mathcal{L}_{\text{CE}}(\phi)$
 - 5: **end for**
 - 6: *Optimize likelihood model:*
 - 7: **for** each gradient step **do**
 - 8: Update $\theta \leftarrow \theta - \lambda_\theta \nabla_\theta \mathcal{L}_{\text{WCE}}(\theta)$
 - 9: Update target parameters: $\bar{\theta} \leftarrow (1 - \alpha)\bar{\theta} + \alpha\theta$
 - 10: **end for**
 - 11: *At inference time, policy probabilities become: $\hat{\pi}(a \mid s) \propto \pi_\phi(a \mid s) \exp(\beta p_\theta(a \mid s))$*
-

Experiment Settings



ALFWorld

Observation:

Your task is to: *put some vase in safe.*
You are in the middle of a room.
Looking quickly around you, you see a drawer 2, a shelf 5, a drawer 1, a shelf 4, a sidetable 1, a drawer 5, a shelf 6,

...

Agent: go to shelf 6

Observation:

You arrive at loc 4. On the shelf 6, you see a vase 2.

Agent: take vase 2 from shelf 6

Figure 2: Overview of all the evaluated tasks, spanning both text and image inputs. Solving all the tasks effectively requires our algorithm to be able to be used to fine-tune LLMs, VLMs, and even robotics transformer models.

Experimental Results

Method	language games			alfworld					
	Chess	Wordle	20Q	Pick	Examine	Clean	Heat	Cool	Pick2
ReAct	0	-4.96	-13.2	45	19	17	7	12	24
SFT	0.11	-3.81	-17.3	38	15	0	11	0	18
ILQL	0.09	-2.08	-14.2	28	7	0	5	2	15
Q-SFT (ours)	0.15	-2.11	-13.1	39	21	19	14	18	21

Table 1: Average scores (for language games), and success rates (for ALFWorld tasks) across 100 independent evaluations. Our method performs best or near-best across the table, and competitively with prompting a much more complex model.

- ReAct: CoT/prompt-based reasoning
- SFT: just using π_ϕ
- ILQL (Implicit Language Q-Learning): train an additional transformer to predict the Q-values directly.

Experimental Results

Method	Score
ReAct	0.60
SFT	0.55
Offline ArCHer	0.57
Q-SFT	0.63

Table 2: Average score across 100 held-out instructions in WebShop. Our method performs best, even against prompting a much larger model.

Method	Pick Object	Place Object Near Target
BC	44	32
CQL	78	57
QT	92	68
Q-SFT	94	64

Table 3: Success rate for 100 runs across robotic manipulation tasks. Our general method performs competitively with Q-transformer, a value-based RL method specifically designed for continuous control.

- ReAct: CoT/prompt-based reasoning
- SFT/BC (behavioural cloning): just using π_ϕ
- Offline ArCHer: hierarchical value modelling at multi-turn-level & token-level (seems complicated).
- CQL (Conservative Q-Learning) & QT (Q-Transformer): train Q-value networks.

Experimental Results




Figure 3: Success rate during initial training on the pick object task of the robotic manipulation benchmark. Though our method achieves similar final performance as Q-transformer, we perform much better on fewer samples.




Figure 4: Scores after training on 10% of the offline dataset on the 20Q task, varying the size of the pretrained model. Our method benefits more from using more sophisticated pretrained models, suggesting our approach scales better.

- (Left) Q-SFT scales better than QT with fewer samples.
- (Right) Q-SFT scales better than ILQL with more parameters. (“We only train on 10% of the dataset, so that retaining prior knowledge from pretraining becomes crucial”...)

Summary of Contributions

- Reframes Q-learning as a supervised fine-tuning problem using a weighted cross-entropy loss (good for stability and simplicity).
- An effective way to leverage pretrained models without adding new layers or heads (this seems worthwhile).
- But requires training two models, π_ϕ and \hat{p}_θ , plus the moving average \bar{p}_θ .
- Some of the choices seem a bit arbitrary (e.g. the moving average to help with stability) and/or not well-motivated/explained.
- No comparison to online methods?